

```

[cloudera@quickstart ~]$ ls
cloudera-manager  Downloads                kerberos  Pictures
workspace
cm_api.py         eclipse                  lib        Public
Desktop          enterprise-deployment.json Music      Templates
Documents        express-deployment.json  parcels    Videos
[cloudera@quickstart ~]$ cd workspace
[cloudera@quickstart workspace]$ ls
[cloudera@quickstart WordCount]$ curl http://www.gutenberg.org/
files/100/100-0.txt | perl -pe 's/^\xEF\xBB\xBF//' > 100-0.txt
% Total    % Received % Xferd  Average Speed   Time    Time
Time Current                        Dload  Upload  Total  Spent
Left  Speed
100 5622k  100 5622k    0    0  390k      0  0:00:14  0:00:14
--:--:--  742k
[cloudera@quickstart WordCount]$ ls
100-0.txt bin conf hadoop.log output src
[cloudera@quickstart WordCount]$ cd ..
[cloudera@quickstart workspace]$ cd ..
[cloudera@quickstart ~]$ pwd
/home/cloudera
[cloudera@quickstart ~]$ hadoop fs -put workspace/WordCount/
100-0.txt
[cloudera@quickstart ~]$ hadoop jar WordCount.jar
edu.stanford.cs246.wordcount.WordCount 100-0.txt output
[100-0.txt, output]
[100-0.txt, output]
21/03/09 05:15:24 INFO client.RMProxy: Connecting to ResourceManager
at /0.0.0.0:8032
21/03/09 05:15:25 INFO input.FileInputFormat: Total input paths to
process : 1
21/03/09 05:15:25 INFO mapreduce.JobSubmitter: number of splits:1
21/03/09 05:15:25 INFO mapreduce.JobSubmitter: Submitting tokens for
job: job_1615290818956_0001
21/03/09 05:15:25 INFO impl.YarnClientImpl: Submitted application
application_1615290818956_0001
21/03/09 05:15:26 INFO mapreduce.Job: The url to track the job:
http://quickstart.cloudera:8088/proxy/
application_1615290818956_0001/
21/03/09 05:15:26 INFO mapreduce.Job: Running job:
job_1615290818956_0001
21/03/09 05:15:35 INFO mapreduce.Job: Job job_1615290818956_0001
running in uber mode : false
21/03/09 05:15:35 INFO mapreduce.Job:  map 0% reduce 0%
21/03/09 05:15:43 INFO mapreduce.Job:  map 100% reduce 0%
21/03/09 05:15:52 INFO mapreduce.Job:  map 100% reduce 100%
21/03/09 05:15:52 INFO mapreduce.Job: Job job_1615290818956_0001
completed successfully
21/03/09 05:15:52 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=11671983
        FILE: Number of bytes written=23567909
        FILE: Number of read operations=0

```

```

FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=5757232
HDFS: Number of bytes written=829099
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots
(ms)=6158
  Total time spent by all reduces in occupied slots
(ms)=6359
  Total time spent by all map tasks (ms)=6158
  Total time spent by all reduce tasks (ms)=6359
  Total vcore-seconds taken by all map tasks=6158
  Total vcore-seconds taken by all reduce tasks=6359
  Total megabyte-seconds taken by all map
tasks=6305792
  Total megabyte-seconds taken by all reduce
tasks=6511616
Map-Reduce Framework
  Map input records=170607
  Map output records=1044475
  Map output bytes=9583027
  Map output materialized bytes=11671983
  Input split bytes=120
  Combine input records=0
  Combine output records=0
  Reduce input groups=76517
  Reduce shuffle bytes=11671983
  Reduce input records=1044475
  Reduce output records=76517
  Spilled Records=2088950
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=133
  CPU time spent (ms)=4360
  Physical memory (bytes) snapshot=361603072
  Virtual memory (bytes) snapshot=3008335872
  Total committed heap usage (bytes)=226365440
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=5757112
File Output Format Counters

```

```
Bytes Written=829099
[cloudera@quickstart ~]$ hadoop fs -ls output
Found 2 items
-rw-r--r--  1 cloudera cloudera          0 2021-03-09 05:15 output/
_SUCCESS
-rw-r--r--  1 cloudera cloudera    829099 2021-03-09 05:15 output/
part-r-00000
[cloudera@quickstart ~]$ hadoop fs -cat output/part\* | head
      82975
"      49
"Air,"  1
"Defects,"  1
"Information  1
"Neither1
"Plain  2
"Project5
"Right  1
"so  1
cat: Unable to write to output stream.
[cloudera@quickstart ~]$
```